

Neural Networks for AI

Lab 3

Fernanda Marana (s3902064)

Floris Cornel (s2724685)

May 16, 2019

1 MLP

1.1 Is it guaranteed that the network finds a solution? Why so?

No, in most cases the network can find the solution, but sometimes - depending on the starting point - it can hit a plateau and never find a solution in which the error is nearly zero.

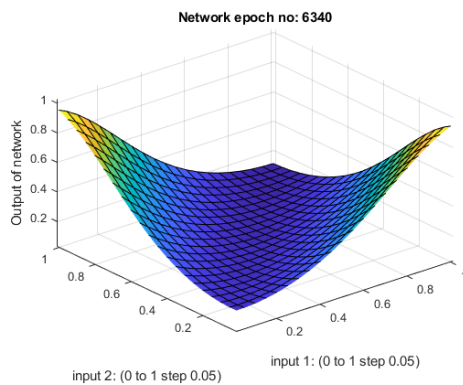


Figure 1: plot




Figure 2: error

1.2 How many epochs are needed to find a solution?

There is no exact value of how many epochs. The network depending on the starting point can take more or less than 5000 epochs to find a solution in which the error is smaller than 0.01.




Figure 3: plot




Figure 4: error

1.3 Set the noise level to 0.5? Explain what happens.

The neural networks - when the noise level is to 0.5 - does not find a solution. This happens because too much noise makes the mapping function too challenging to learn.




Figure 5: plot




Figure 6: error

1.4 Set the noise level to 0.2. Change the weight spread to 5. What can you observe? Explain your results using the delta-rule..

With the noise still being big to confuse the network on the learning process and by having more distance between the weights, the network does not find the ideal goal although it finds values in which the total error becomes less than 0.01.




Figure 7: plot

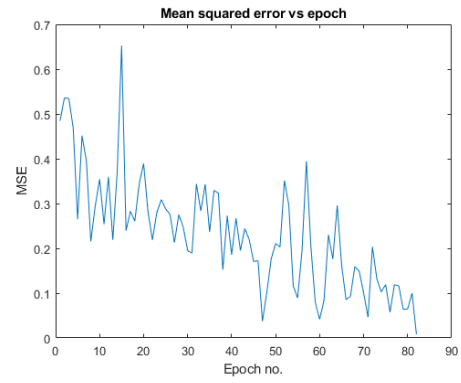


Figure 8: error

1.5 Set the noise level to 0.01. Leave the weight spread at 5. There are two qualitatively different solutions to the XOR problem. What are these two? Include a figure of both solutions.

The two qualitatively different solutions to the XOR problems. One of them finds the goal result very quickly and the other other skips it and never finds this. This is because: 1) The noise is so small it barely has any effect on reducing the generalization error and 2) the weight spread is too bigger which is a problem for the initialization and may be responsible for causing the vanishing gradient situation on the neural network.




Figure 9: The first solution.




Figure 10: The second solution.




Figure 11: Network did not find the solution

1.6 Which shape does the graph of the error usually have? Explain the shape with the use of the delta- rule.

The shapes that the graph usually have is of an hyperbolic paraboloid since it the network tries to approximate to the goal values - and thus minimize the error- by applying a gradient descendant learning rules in which the weights are updated .

2 Another function

2.1 Is the network capable of learning the sine function?

As you can see in the figures below, the network is capable of learning the sin function.




Figure 12: Plot of learned sine




Figure 13: The corresponding error plot

2.2 Set n examples in the top of the file to 5. Rerun the simulation. What can you observe? With which feature of neural networks does this phenomenon correspond?

With less sample points of the sin, the representation of the sin is still very accurate.

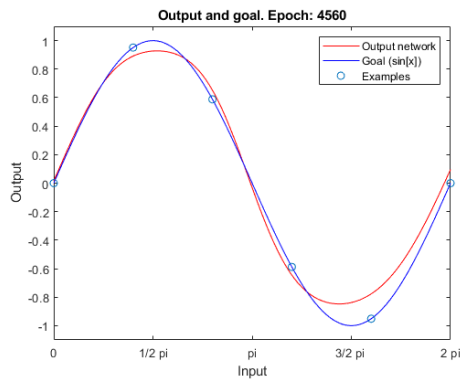


Figure 14: Learned sine with 5 example points

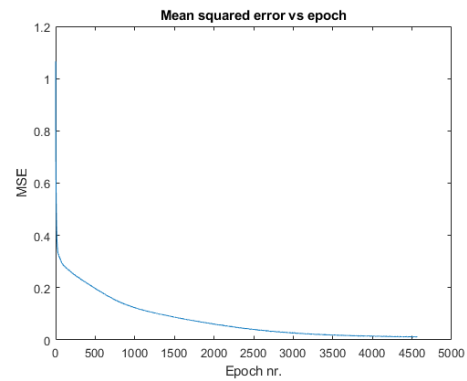


Figure 15: Error of learned sine with 5 example points

2.3 Set plot bigger picture to true. How is the domain of the network determined? What happens if the input is outside of this domain?

The domain of the pi goes from 0 to 2π . Outside of the domain, we can see that the output quickly becomes less accurate.

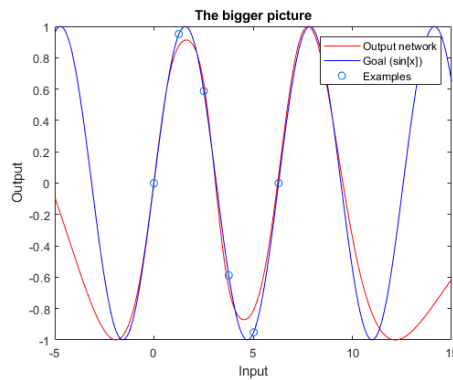


Figure 16: Extended plot of sine

2.4 At least how many neurons are required to learn a sine?

After a lot of simulations, we were able to learn the sine with 3 hidden neurons. It was impossible for us to learn the sine with only 2 hidden neurons.




Figure 17: Plotting sine with 3 hidden neurons




Figure 18: Error sine with 3 hidden neurons

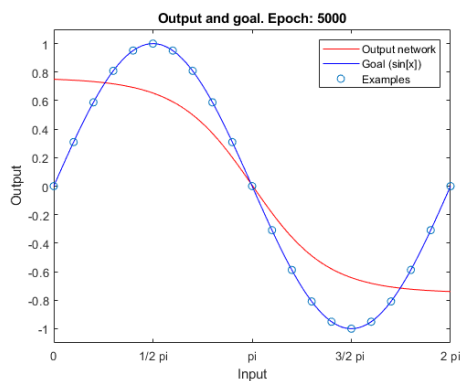


Figure 19: Plotting sine with 2 hidden neurons




Figure 20: Error sine with 2 hidden neurons

2.5 You have modified the output function for this part of the lab assignment. Does the XOR learning network still work? Why so?

We had to modify the output function for this assignment in order to correspond with the newly set goals which are both based on the sin function.

```
1 function [output] = output_function(x)
2     % set output here
3     output = sin(x);
4 end
```

Listing 1: output_function.m

3 Code

```
1 % mlp.m Implementation of the Multi-Layer Perceptron
2
3 clear all
4 close all
5
6 examples = [0 0;1 0;0 1;1 1];
7 goal = [0.01 0.99 0.99 0.01]';
```

```

8
9 % Boolean for plotting the animation
10 plot_animation = true;
11
12
13 % Parameters for the network
14 learn_rate = 0.2;           % learning rate
15 max_epoch = 5000;          % maximum number of epochs
16
17 mean_weight = 0;
18 weight_spread = 5;
19
20 n_input = size(examples,2);
21 n_hidden = 20;
22 n_output = size(goal,2);
23
24 % Noise level at the input
25 noise_level = 0.01;
26
27 % Activation of the bias node
28 bias_value = -1;
29
30
31 % Initializing the weights
32 w_hidden = rand(n_input + 1, n_hidden) .* weight_spread - weight_spread/2 + mean_weight;
33 w_output = rand(n_hidden, n_output) .* weight_spread - weight_spread/2 + mean_weight;
34
35 % Start training
36 stop_criterium = 0;
37 epoch = 0;
38 min_error = 0.01;
39
40 while ~stop_criterium
41     epoch = epoch + 1;
42
43     % Add noise to the input data.
44     noise = randn(size(examples)) .* noise_level;
45     input_data = examples + noise;
46
47     % Append bias to input data
48     input_data(:,n_input+1) = ones(size(examples,1),1) .* bias_value;
49
50     epoch_error = 0;
51     epoch_delta_hidden = 0;
52     epoch_delta_output = 0;
53
54
55
56 % FROM HEREON YOU NEED TO MODIFY THE CODE!
57 for pattern = 1:size(input_data,1)
58
59     % Compute the activation in the hidden layer
60     hidden_activation = input_data(pattern,:) * w_hidden ;

```

```

61
62     % Compute the output of the hidden layer (don't modify this)
63     hidden_output = sigmoid(hidden_activation);
64
65     % Compute the activation of the output neurons
66     %not sure
67     output_activation = hidden_output * w_output ;
68
69     % Compute the output
70     output = output_function(output_activation);
71
72     % Compute the error on the output
73     output_error = (output - goal(pattern));
74
75     % Compute local gradient of output layer
76     local_gradient_output = d_sigmoid(output_activation).*(goal(pattern) -output);
77
78
79     % Compute the error on the hidden layer (backpropagate)
80     hidden_error = 0 ;
81
82     % Compute local gradient of hidden layer
83     local_gradient_hidden = d_output_function(hidden_activation) .* (
84     local_gradient_output * transpose( w_output));
85
86     % Compute the delta rule for the output
87     delta_output = learn_rate * transpose( hidden_output) * local_gradient_output ;
88
89     % Compute the delta rule for the hidden units;
90     delta_hidden = learn_rate * transpose(input_data(pattern,:)) *
91     local_gradient_hidden ;
92
93     % Update the weight matrices
94     w_hidden = upW(w_hidden, delta_hidden);
95     w_output = upW(w_output , delta_output);
96
97     % Store data
98     epoch_error = epoch_error + (output_error).^2;
99     epoch_delta_output = epoch_delta_output + sum(sum(abs(delta_output)));
100     epoch_delta_hidden = epoch_delta_hidden + sum(sum(abs(delta_hidden)));
101
102     end
103
104
105     % Log data
106
107     h_error(epoch) = sum(epoch_error) / size(input_data,1);
108     log_delta_output(epoch) = epoch_delta_output;
109     log_delta_hidden(epoch) = epoch_delta_hidden;
110
111     % Check whether maximum number of epochs is reached

```



```

112 % Implement a stop criterion here
113 if h_error(epoch) < min_error || epoch == max_epoch
114
115     stop_criterium = 1;
116 end
117
118 % Plot the animation
119 if and((mod(epoch,20)==0),(plot_animation))
120     emp_output = zeros(21,21);
121     figure(1)
122     for x1 = 1:21
123         for x2 = 1:21
124             hidden_act = sigmoid([(x1/20 - 0.05) (x2/20 -0.05) bias_value] *
w_hidden);
125             emp_output(x1,x2) = output_function(hidden_act * w_output);
126         end
127     end
128     surf(0:0.05:1,0:0.05:1,emp_output)
129     title(['Network epoch no: ' num2str(epoch)]);
130     xlabel('input 1: (0 to 1 step 0.05)')
131     ylabel('input 2: (0 to 1 step 0.05)')
132     zlabel('Output of network')
133     zlim([0 1])
134     pause(0.01)
135 end
136
137 end
138
139 % Plotting the error
140 figure(2)
141 plot(1:epoch,h_error)
142 title('Mean squared error vs epoch');
143 xlabel('Epoch no. ');
144 ylabel('MSE');
145
146 % Add additional plot functions here (optional)

```

Listing 2: mlp.m

```

1 clear all
2 close all
3
4 % The number of examples taken from the function
5 n_examples = 5;
6
7 examples = (0:2*pi/(n_examples):2*pi).';
8 goal = sin(examples);
9
10 % Boolean for plotting animation
11 plot_animation = true;
12 plot_bigger_picture = true;
13
14 % Parameters for the network
15 learn_rate = 0.05 ; % learning rate

```

```

16 max_epoch = 5000; % maximum number of epochs
17
18
19 mean_weight = 0;
20 weight_spread = 1;
21
22 n_input = size(examples,2);
23 n_hidden = 50;
24 n_output = size(goal,2);
25
26 % Noise level at input
27 noise_level = 0.01;
28
29 bias_value = -1;
30
31 % Initializing the weights
32 w_hidden = rand(n_input + 1, n_hidden) .* weight_spread - weight_spread/2 + mean_weight;
33 w_output = rand(n_hidden, n_output) .* weight_spread - weight_spread/2 + mean_weight;
34
35 % Start training
36 stop_criterium = 0;
37 epoch = 0;
38 min_error = 0.01;
39
40 while ~stop_criterium
41     epoch = epoch + 1;
42
43     % Add noise to the input
44     noise = randn(size(examples)) .* noise_level;
45     input_data = examples + noise;
46
47     % Append bias
48     input_data(:,n_input+1) = ones(size(examples,1),1) .* bias_value;
49
50     epoch_error = 0;
51     epoch_delta_hidden = 0;
52     epoch_delta_output = 0;
53     for pattern = 1:size(input_data,1)
54         % Compute the activation in the hidden layer
55         hidden_activation = input_data(pattern,:) * w_hidden;
56
57         % Compute the output of the hidden layer (don't modify this)
58         hidden_output = sigmoid(hidden_activation);
59
60         % Compute the activation of the output neurons
61         %not sure
62         output_activation = hidden_output * w_output;
63
64         % Compute the output
65         output = output_function(output_activation);
66
67         % Compute the error on the output
68         output_error = (output - goal(pattern));

```

```

69
70     % Compute local gradient of output layer
71     local_gradient_output = d_sigmoid(output_activation).*(goal(pattern) -output);
72
73
74     % Compute the error on the hidden layer (backpropagate)
75     hidden_error = 0 ;
76
77     % Compute local gradient of hidden layer
78     local_gradient_hidden = d_output_function_sin(hidden_activation) .* (
79     local_gradient_output * transpose( w_output));
80
81     % Compute the delta rule for the output
82     delta_output = learn_rate * transpose( hidden_output) * local_gradient_output ;
83
84     % Compute the delta rule for the hidden units;
85     delta_hidden = learn_rate * transpose(input_data(pattern,:)) *
86     local_gradient_hidden ;
87
88     % Update the weight matrices
89     w_hidden = upW(w_hidden, delta_hidden);
90     w_output = upW(w_output , delta_output);
91
92     % Store data
93     epoch_error = epoch_error + (output_error).^2;
94     epoch_delta_output = epoch_delta_output + sum(sum(abs(delta_output)));
95     epoch_delta_hidden = epoch_delta_hidden + sum(sum(abs(delta_hidden)));
96
97
98     end
99
100     h_error(epoch) = sum(epoch_error) / size(input_data,1);
101     log_delta_output(epoch) = epoch_delta_output;
102     log_delta_hidden(epoch) = epoch_delta_hidden;
103
104     if epoch > max_epoch
105         %stop_criterium = 1;
106     end
107
108     % Add your stop criterion here
109     if h_error(epoch) < min_error || epoch == max_epoch
110
111         stop_criterium = 1;
112     end
113
114     % Plot the animation
115     if and((mod(epoch,20)==0),(plot_animation))
116         %out = zeros(21,1);
117         nPoints = 100;
118         input = linspace(0, 2 * pi, nPoints);
119         for x=1:nPoints
120             h_out = sigmoid([input(x) bias_value] * w_hidden);
121             out(x) = output_function(h_out * w_output);
122         end
123     end

```

```

120     end
121     figure(1)
122     plot(input,out,'r-','DisplayName','Output network')
123     hold on
124     plot(input,sin(input),'b-','DisplayName','Goal (sin[x])')
125     hold on
126     scatter(examples, goal, 'DisplayName', 'Examples')
127     hold on
128     title(['Output and goal. Epoch: ' num2str(epoch)]);
129     xlim([0 2*pi])
130     ylim([-1.1 1.1])
131     set(gca,'XTick',0:pi/2:2*pi)
132     set(gca,'XTickLabel',{'0','1/2 pi','pi','3/2 pi ','2 pi'})
133     xlabel('Input')
134     ylabel('Output')
135     legend('location','NorthEast')
136     hold off
137     pause(0.01)
138 end
139
140 end
141
142
143 % Plot error
144 figure(2)
145 plot(1:epoch,h_error)
146 title('Mean squared error vs epoch');
147 xlabel('Epoch nr. ');
148 ylabel('MSE');
149
150 %Plot the bigger picture
151 if plot_bigger_picture
152     figure(3)
153     in_raw = (-5:0.1:15)';
154     in_raw = horzcat(in_raw,(bias_value*ones(size(in_raw))));
155     h_big = sigmoid(in_raw * w_hidden);
156     o_big = output_function(h_big * w_output);
157
158     plot(-5:0.1:15,o_big,'r-','DisplayName','Output network')
159     hold on
160     plot(-5:0.1:15,sin(-5:0.1:15),'b-','DisplayName','Goal (sin[x])')
161     hold on
162     scatter(examples, sin(examples), 'DisplayName', 'Examples');
163     hold off
164     xlabel('Input')
165     ylabel('Output')
166     legend('location','NorthEast')
167     title('The bigger picture')
168 end

```

Listing 3: mlp_sinus.m