

Neural Networks for AI

LAB 4

Fernanda Marana (s3902064)

Floris Cornel (s2724685)

May 23, 2019

1 Questions on the Hopfield implementation




Figure 1: Weight matrix



Figure 2: Correct recognition of 3 patterns.

In figure 1 you can see the normalized weight matrix. Figure 2 depicts a correct output found by the network.

1.1 Explain why repetitively applying Hebb's learning rule can be captured by the following formula

Hebb's learning rule is based on the idea that nodes that tend to have the same value over the training set are strongly correlated, and nodes with opposite values are anti-correlated. The correlated nodes get more positive values and the anti-correlated nodes get more negative values. Doing this process repetitively, can be done by looking at v_i^p and v_i^p for every pattern p . The weights will be determined by the sum of $-1 * -1 = 1$ and $1 * 1 = 1$ for correlated nodes, and $-1 * 1 = -1$ and $1 * -1 = -1$ for anti-correlated nodes. So we can derive $w_{ij} = \sum_{p=1}^P v_i^p * v_j^p$.

- 1.2 Omit the normalization of the weights and rerun the algorithm. What happens to the performance of the network? How can you explain this?**




Figure 3: Weight matrix of without normalization




Figure 4: Correct recognition of 3 patterns without normalization.

What we can see from figures 3 and 4 is that when the normalization is turned off, the range of values in the weight matrix goes from 3 to -3 instead of from 1 to -1. However, it does not affect the networks ability to recognise the three patterns. This is because the weights, whether normalized or not, affect the output relatively to to other weights, so normalisation is not strictly needed.

- 1.3 What is the theoretical number of patterns that a Hopfield network using 25 neurons can store?**

Following the equation bellow

$$\frac{N}{2 * \ln(N)}$$

in which N corresponds to the number of neurons, we get the value of 3.88. So the theoretical number of patterns must be below it.

- 1.4 Vary $n_examples$. How many different patterns can be stored properly by the network? Why is this different from what you have found in (1.3)?**

We were unable to successfully recognise 4 different patterns with our network, confirming the $3 < 3.88 < 4$ as the maximum number of stored patterns.

2 Noise and spurious states



Figure 5: Output with deviant end states




Figure 6: Output of inverted

2.1 Explain using your own words how the network is able to recover the patterns from the noisy input. Your answer should at least involve energy and weights.

The network is able to recover the patterns from the noisy because the energy is responsible to shape the network either decreasing or staying the same upon units being updated. So under repeated updating the network will eventually converge to a state which is a local minimum in the energy function. Additionally, the weights guarantees that the energy function decreases monotonically while following the activation rules.

2.2 Occasionally, for higher noise levels (30 - 35%) some peculiar end states might occur. Explain why these states are only observed with synchronous updates and not with asynchronous updates. If you cannot find any deviant end states then you can increase the chance of observing one by using more patterns and increasing the amount of epochs.

In figure 5 we see an example of output with deviant end states due to higher noise. This mismatched recognition is because the flipping of random bits makes the input more similar to a wrong goal than its own correct goal. It does not always happen, but the larger the noise level, the larger the chance of becoming similar to the wrong goal. Very low noise levels makes it impossible because there is a minimum number of flips that need to take place in order to cross the line between one goal and another goal.

2.3 Set invert in the beginning of the code to true. Explain what happens to the input. Does anything change for the training of the network?

We can see in figure 6 that the input pattern is inverted, but the network is still capable of learning the correct shape, so the learning is not effected negatively.

2.4 Explain why the inverted patterns are also stored by the network.

Well, the network looks at the edges of certain patterns, so the difference between certain neighbours. These differences can either be large or small, but there is no indication whether a large difference should be positive or negative. As a result, the inverted pattern is just as valid and easily recognised as the regular pattern.

3 Code

```
1 % Clear workspace and close all previous windows
2 clear all;
3 close all;
4
5 % Initializing data and parameters
6 % PARAMETERS
7 n_examples = 3;           % The number of examples(0 < n_examples < 7)
8 n_epochs = 6;             % The number of epochs
9 normalize_weights = true; % Normalization bool
10
11 random_percentage = 60;    % Percentage of bits that are flipped randomly
12 invert = false;           % Invert the input (test for spurious states)
13
14
15 % Do not change these lines
16 dim_x = 5;                % Dimensions of examples
17 dim_y = 5;
18
19 % Compute size of examples
20 size_examples = dim_x * dim_y;
21
22 % Convert percentage to fraction
23 random_percentage = random_percentage/100;
24
25 % Set color for network plots
26 color = 20;
27
28 % The data is stored in .dat files. They have to be located in the same
29 % directory as this source file
30 data = importdata('M.dat');
31 data(:, :, 2) = importdata('A.dat');
32 data(:, :, 3) = importdata('S.dat');
33 data(:, :, 4) = importdata('T.dat');
34 data(:, :, 5) = importdata('E.dat');
35 data(:, :, 6) = importdata('R.dat');
36
37 % Convert data matrices into row vectors. Store all vectors in a matrix
38 vector_data = zeros(n_examples, size_examples);
39 for idx = 1:n_examples
40     % Every row will represent an example
41     vector_data(idx, :) = reshape(data(:, :, idx)', 1, size_examples);
42 end
43
44
```

```

45 % TRAINING THE NETWORK
46 % The network is trained using one-shot Hebbian learning
47
48 % The result should be a matrix dimensions: size_examples * size_examples
49 % Each entry should contain a sum of n_examples
50 weights = transpose(vector_data) * vector_data;
51
52 % A hopfield neuron is not connected to itself. The diagonal of the matrix
53 % should be zero.
54 weights = weights - diag(diag(weights));
55
56 % These lines check whether the matrix is a valid weight matrix for a
57 % Hopfield network.
58 assert(isequal(size(weights),[size_examples size_examples]), ...
59         'The matrix dimensions are invalid');
60 assert(isequal(tril(weights)',triu(weights)), ...
61         'The matrix is not symmetric');
62 assert isempty(find(diag(weights), 1)), ...
63         'Some neurons are connected to themselves');
64
65 % Normalizing the weights
66 if normalize_weights
67     weights = weights ./ n_examples;
68 end
69
70
71 % PLOT WEIGHT MATRIX
72 figure(1)
73 imagesc(weights)
74 colorbar
75 title('Weight matrix')
76
77 % INTRODUCE NOISE
78
79 % Copy the input data
80 input = vector_data;
81
82 % Create a matrix with the same dimensions as the input data in which
83 % random_percentage elements are set to -1 and the others are set to 1.
84 % We do this by sampling from a normal distribution
85 noise_matrix = (randn(size(input)) > norminv(random_percentage));
86 noise_matrix = noise_matrix - (noise_matrix==0);
87
88 % Flip bits (* -1) using the noise_matrix
89 input = input .* noise_matrix;
90
91 % Optionally invert the input
92 if invert
93     input = -1 .* input;
94 end
95
96 % PLOTTING INPUT PATTERNS
97 figure(2)

```

```

98 for example = 1:n_examples
99     subplot(n_epochs + 2,n_examples,example)
100     test = reshape((input(example,:)),dim_x, dim_y)';
101     image(test .* color + color)
102     str = 'Ex: ';
103     str = strcat(str,int2str(example));
104     title(str)
105     if(example == 1)
106         axis on
107         ylabel('input')
108     else
109         axis off
110     end
111     axis square
112 end
113
114 % UPDATING THE NETWORK
115 % Feed the network with all of the acquired inputs. Update the network and
116 % plot the activation after each epoch.
117 for example = 1:n_examples
118
119     % The initial activation is the row vector of the current example.
120     activation = input(example,:)';
121
122     for epoch = 1:n_epochs
123         % Compute the new activation
124         activation = weights * activation;
125
126         % Apply the activation function
127         for i=1:size(activation)
128             if activation(i)>=0
129                 activation(i) = 1;
130             else
131                 activation(i) = 0;
132             end
133         end
134
135         % PLOTTING THE ACTIVATION
136
137         % Reshape the activation such that we get a 5x5 matrix
138         output = reshape(activation, dim_x, dim_y)';
139
140         % Compute the index of where to plot
141         idx = epoch * n_examples + example;
142
143         % Create the plot
144         subplot(n_epochs + 2,n_examples,idx)
145         image(output .* color + epoch * color)
146
147         % Only draw axes on the leftmost column
148         if(example == 1)
149             axis on
150             str = 'Ep: ';

```

```

151         str = strcat(str,int2str(epoch));
152         ylabel(str)
153     else
154         axis off
155     end
156
157     % Make sure the plots use a square grid
158     axis square
159 end
160 end
161
162 % Finally we plot the goal vector for comparison
163 for idx = 1:n_examples
164     subplot(n_epochs + 2,n_examples,(n_epochs + 1) * n_examples + idx);
165     image(data(:,:,idx).* color + n_epochs+1 * color)
166     if(idx == 1)
167         axis on
168         ylabel('goal')
169     else
170         axis off
171     end
172     axis square
173 end
174 %-----

```

Listing 1: hopfield.m