

Neural Networks for AI

LAB 5

Fernanda Marana (s3902064)

Floris Cornel (s2724685)

June 6, 2019

1 Experiments

- 1.1 Convince the reader of your report that you have understood the network. Mention all parameters that are relevant to the training procedure and what their role is. This includes the parameters outside your own implemented pieces of code. Do not explain formulas/loopsetc. Make sure your experiment is reproducible by anybody that at least understands CNNs.**

As explained on the laboratory document, the main purpose of this exercise was to implement a CNN to classify images of handwritten digits in the correct prediction. The creation of a CNN - like a typical neural network - also consists on inputs, hidden layers and outputs, but they already assume that the input are images and, as a consequence, the layers have neurons arranged in 3 dimensions: width, height, depth. Additionally, since images consist on great amount of volumes, CNNs are known for having three very important processes: the convolution, the pooling and the normalization (loss) function.

The convolution consists on going through each image and each filter so the activation map can be created. Also, on this step the network starts to learn filters that activate when they detect some type of visual feature. So, by doing the `cnnConvolve` function, the learning process begins with the specific parameters set on the network. An important note is that the weights are rotated two times to do the convolution (otherwise, we had be doing the cross- correlation, i.e., measuring the similarity between two signals, which is not the objective). The forward propagation in this exercise does the activation of neurons with the convolution, pooling and softmax function while the backpropagation receives the error obtained by the result and backpropagates it through the softmax and subsampling layers.

Pooling's purpose is to reduce the dimensions of the data by obtaining an average, in this case, between a defined number of elements on the matrix. Since images can have big dimensions and we can be dealing with a great number of them, reducing the spatial size of the representation can help the amount of computation in the network be also reduced.

The softmax function is one type of the normalization layers that is responsible to measure the compatibility between a prediction and the ground truth label. In other words, the prediction would be what the neural network detect as the number that is currently processing and the truth label would be the actually number that is represented by the handwriting.

Finally, some important point worthy noticing is that the forward propagation and back propagation are still used to receive - from the input given- a response and "reshape" (do some alterations) on the network based on it.

1.2 What can you say about the filters now when compared to the beginning of the training phase? Explain.

It is possible to see that the filters are starting to create a pattern in which they identify a feature of the image. In the beginning, they were very random but as the network start to learn each number, the patterns are being created. As explained on the section 1.1, the reason why this is happening is due to the forward and back propagation on the network that are helping the learning of what filters activate when they detect some type of visual feature.




Figure 1: Weight matrix




Figure 2: Correct recognition of 3 patterns.

1.3 Report the accuracy of the network on the test set after 3 epochs. It should be higher than 97 % .

The accuracy reached is 0.971600

2 Code

```
1 function convolvedFeatures = cnnConvolve(filterDim, numFilters, images, W, b)
2 %cnnConvolve Returns the convolution of the features given by W and b with
3 %the given images
4 %
5 % Parameters:
6 % filterDim - filter (feature) dimension
7 % numFilters - number of feature maps
8 % images - large images to convolve with, matrix in the form
9 %           images(r, c, image number)
10 % W, b - W, b for features from the sparse autoencoder
11 %       W is of shape (filterDim,filterDim,numFilters)
12 %       b is of shape (numFilters,1)
13 %
14 % Returns:
15 % convolvedFeatures - matrix of convolved features in the form
16 %                   convolvedFeatures(imageRow, imageCol, featureNum, imageNum)
17
18 numImages = size(images, 3);
19 imageDim = size(images, 1);
20 convDim = imageDim - filterDim + 1;
21
```

```

22 convolvedFeatures = zeros(convDim, convDim, numFilters, numImages);
23
24 % Instructions:
25 %   Convolve every filter with every image here to produce the
26 %   (imageDim - filterDim + 1) x (imageDim - filterDim + 1) x numFeatures x numImages
27 %   matrix convolvedFeatures, such that
28 %   convolvedFeatures(imageRow, imageCol, featureNum, imageNum) is the
29 %   value of the convolved featureNum feature for the imageNum image over
30 %   the region (imageRow, imageCol) to (imageRow + filterDim - 1, imageCol + filterDim -
31 %   1)
32 %
33 % Expected running times:
34 %   Convolving with 100 images should take less than 30 seconds
35 %   Convolving with 5000 images should take around 2 minutes
36 %   (So to save time when testing, you should convolve with less images, as
37 %   described earlier)
38 %%% Add code here
39
40 for k=1:numImages
41     im = images(:,:,k);
42     for l=1:numFilters
43         weight = rot90(rot90(W(:,:,l)));
44         result = conv2(im,weight,'valid');
45
46         convolvedFeatures(:,:,l,k) = logsig(b(l)+result);
47     end
48 end
49
50
51
52 end

```

Listing 1: cnnConvolve.m

```

1 function pooledFeatures = cnnPool(poolDim, convolvedFeatures)
2 %cnnPool Pools the given convolved features
3 %
4 % Parameters:
5 %   poolDim - dimension of pooling region
6 %   convolvedFeatures - convolved features to pool (as given by cnnConvolve)
7 %                       convolvedFeatures(imageRow, imageCol, featureNum, imageNum)
8 %
9 % Returns:
10 %   pooledFeatures - matrix of pooled features in the form
11 %                   pooledFeatures(poolRow, poolCol, featureNum, imageNum)
12 %
13
14 numImages = size(convolvedFeatures, 4);
15 numFilters = size(convolvedFeatures, 3);
16 convolvedDim = size(convolvedFeatures, 1);
17
18 pooledFeatures = zeros(convolvedDim / poolDim, ...
19     convolvedDim / poolDim, numFilters, numImages);

```

```

20
21
22
23 % Instructions:
24 %   Now pool the convolved features in regions of poolDim x poolDim,
25 %   to obtain the
26 %   (convolvedDim/poolDim) x (convolvedDim/poolDim) x numFeatures x numImages
27 %   matrix pooledFeatures, such that
28 %   pooledFeatures(poolRow, poolCol, featureNum, imageNum) is the
29 %   value of the featureNum feature for the imageNum image pooled over the
30 %   corresponding (poolRow, poolCol) pooling region.
31 %
32 %   Use mean pooling here.
33
34 %%% Add code here
35 separation = poolDim * ones(1, convolvedDim / poolDim);
36
37 for i=1:numImages
38     for j=1:numFilters
39
40         d=mat2cell(convolvedFeatures(:,:,j,i),separation, separation);
41         new = zeros(size(d,1));
42         for k=1:size(d,1)
43             for r=1:size(d,1);
44                 new(k,r) = mean2(d{[k,r]});
45             end
46         end
47
48         % result = reshape(new, (convolvedDim/poolDim),(convolvedDim/poolDim));
49         pooledFeatures(:,:,j,i) = new;
50     end
51 end
52 end
53
54 end

```

Listing 2: cnnPool

```

1 function [cost, grad, preds, activations] = cnnCost(theta,images,labels,numClasses,...
2             filterDim,numFilters,poolDim,pred)
3 % Calcualte cost and gradient for a single layer convolutional
4 % neural network followed by a softmax layer with cross entropy
5 % objective.
6 %
7 % Parameters:
8 %   theta      - unrolled parameter vector
9 %   images     - stores images in imageDim x imageDim x numImages
10 %               array
11 %   numClasses - number of classes to predict
12 %   filterDim  - dimension of convolutional filter
13 %   numFilters - number of convolutional filters
14 %   poolDim    - dimension of pooling area
15 %   pred       - boolean only forward propagate and return
16 %               predictions

```

```

17 %
18 %
19 % Returns:
20 % cost      - cross entropy cost
21 % grad      - gradient with respect to theta (if pred==False)
22 % preds     - list of predictions for each example (if pred==True)
23
24
25 if ~exist('pred','var')
26     pred = false;
27 end;
28
29
30 imageDim = size(images,1); % height/width of image
31 numImages = size(images,3); % number of images
32
33 %% Reshape parameters and setup gradient matrices
34
35 % Wc is filterDim x filterDim x numFilters parameter matrix
36 % bc is the corresponding bias
37
38 % Wd is numClasses x hiddenSize parameter matrix where hiddenSize
39 % is the number of output units from the convolutional layer
40 % bd is corresponding bias
41 [Wc, Wd, bc, bd] = cnnParamsToStack(theta,imageDim,filterDim,numFilters,...
42                                     poolDim,numClasses);
43
44 % Same sizes as Wc,Wd,bc,bd. Used to hold gradient w.r.t above params.
45 Wc_grad = zeros(size(Wc));
46 Wd_grad = zeros(size(Wd));
47 bc_grad = zeros(size(bc));
48 bd_grad = zeros(size(bd));
49 Wc_grad_size = (size(Wc));
50 Wd_grad_size = size(Wd);
51 bc_grad_size = size(bc);
52 bd_grad_size = size(bd);
53 %=====
54 %% Forward Propagation
55 % In this step you will forward propagate the input through the
56 % convolutional and subsampling (mean pooling) layers. You will then use
57 % the responses from the convolution and pooling layer as the input to a
58 % standard softmax layer.
59
60 %% Convolutional Layer
61 % For each image and each filter, convolve the image with the filter, add
62 % the bias and apply the sigmoid nonlinearity. Then subsample the
63 % convolved activations with mean pooling. Store the results of the
64 % convolution in activations and the results of the pooling in
65 % activationsPooled. You will need to save the convolved activations for
66 % backpropagation.
67 convDim = imageDim-filterDim+1; % dimension of convolved output
68 outputDim = (convDim)/poolDim; % dimension of subsampled output
69

```

```

70 % convDim x convDim x numFilters x numImages tensor for storing activations
71
72 %%% REPLACE THE FOLLOWING LINE %%%
73 activations = zeros(convDim,convDim,numFilters,numImages);
74 activations = cnnConvolve(filterDim, numFilters, images, Wc, bc);
75 % outputDim x outputDim x numFilters x numImages tensor for storing
76 % subsampled activations
77 size (activations);
78 %%% REPLACE THE FOLLOWING LINE %%%
79 activationsPooled = zeros(outputDim,outputDim,numFilters,numImages);
80 activationsPooled = cnnPool(poolDim,activations );
81 size (activationsPooled);
82
83 % Reshape activations into 2-d matrix, hiddenSize x numImages,
84 % for Softmax layer
85 %%% REPLACE THE FOLLOWING LINE %%%
86 activationsPooled = reshape(activationsPooled,size(Wd,2) , numImages);
87
88 %% Softmax Layer
89 % Forward propagate the pooled activations calculated above into a
90 % standard softmax layer. For your convenience we have reshaped
91 % activationPooled into a hiddenSize x numImages matrix. Store the
92 % results in probs.
93
94 % numClasses x numImages for storing probability that each image belongs to
95 % each class.
96
97 %%% COMPUTE THE SOFTMAX OUTPUT %%%
98 probs = zeros(numClasses,numImages);
99 probs = Wd*activationsPooled;
100 for j=1: size(probs,1)
101     probs(j,:) = probs(j,:) + bd(j);
102 end
103
104 Y_num = exp(probs);
105
106
107 for k=1: size(Y_num,2)
108     total_sum = sum(Y_num(:,k));
109     Y_num(:,k)= Y_num(:,k)/total_sum;
110
111 end
112 probs = Y_num;
113 %=====
114 %% STEP 1b: Calculate Cost
115 % In this step you will use the labels given as input and the probs
116 % calculate above to evaluate the cross entropy objective. Store your
117 % results in cost.
118 indexes = sub2ind(size(probs), labels', 1:numImages);
119 cost = -mean(log(probs(indexes)));
120
121 % Makes predictions given probs and returns without backproagating errors.
122 [~,preds] = max(probs,[],1);

```

```

123 preds = preds';
124 if pred
125     grad = 0;
126     return;
127 end;
128
129 %=====
130 %% STEP 1c: Backpropagation
131 % Backpropagate errors through the softmax and convolutional/subsampling
132 % layers. Store the errors for the next step to calculate the gradient.
133 % Backpropagating the error w.r.t the softmax layer is as usual. To
134 % backpropagate through the pooling layer, you will need to upsample the
135 % error with respect to the pooling layer for each filter and each image.
136 % Use the kron function and a matrix of ones to do this upsampling
137 % quickly.
138
139 deriv = probs;
140 deriv(indexes) = deriv(indexes) - 1;
141 deriv = deriv ./ numImages;
142
143 Wd_grad = deriv * activationsPooled';
144 bd_grad = sum(deriv, 2);
145
146 deriv2_pooled = Wd' * deriv;
147 deriv2_pooled = reshape(deriv2_pooled, outputDim, outputDim, numFilters, numImages);
148 delta_upsampled = zeros(convDim, convDim, numFilters, numImages);
149
150 for im_idx=1:numImages
151     im = squeeze(images(:,:,im_idx));
152     for f_idx=1:numFilters
153         delta_pool = (1/poolDim^2) * kron(squeeze(deriv2_pooled(:,:,f_idx,im_idx)), ones
            (poolDim));
154         delta_upsampled(:,:,f_idx, im_idx) = delta_pool .* ...
            activations(:,:,f_idx,im_idx).*(1-activations(:,:,f_idx,im_idx));
155         delta_pool_sqz = squeeze(delta_upsampled(:,:,f_idx,im_idx));
156         cur_grad = conv2(im, rot90(delta_pool_sqz, 2), 'valid');
157
158         Wc_grad(:,:,f_idx) = Wc_grad(:,:,f_idx) + cur_grad;
159         bc_grad(f_idx) = bc_grad(f_idx) + sum(delta_pool_sqz(:));
160     end
161 end
162
163
164 %% Unroll gradient into grad vector for minFunc
165 grad = [Wc_grad(:) ; Wd_grad(:) ; bc_grad(:) ; bd_grad(:)];
166
167 end

```

Listing 3: cnnPool